

02_python_pandas

December 9, 2020

1 Pandas

1.1 Introduction

pandas is the Python Data Analysis package. It allows for data ingestion, transformation and cleaning, and creates objects that can then be passed on to analytic packages like **statsmodels** and **scikit-learn** for modeling and packages like **matplotlib**, **seaborn**, and **plotly** for visualization.

pandas is built on top of **numpy**, so many **numpy** functions are commonly used in manipulating **pandas** objects.

pandas is a pretty extensive package, and we'll only be able to cover some of its features. For more details, there is free online documentation at pandas.pydata.org. You can also look at the book “[Python for Data Analysis \(2nd edition\)](#)” by Wes McKinney, the original developer of the **pandas** package, for more details.

1.2 Starting pandas

As with any Python module, you have to “activate” **pandas** by using **import**. The “standard” alias for **pandas** is **pd**. We will also import **numpy**, since **pandas** uses some **numpy** functions in the workflows.

```
[1]: import numpy as np
import pandas as pd
```

1.3 Data import and export

Most data sets you will work with are set up in tables, so are rectangular in shape. Think Excel spreadsheets. In **pandas** the structure that will hold this kind of data is a **DataFrame**. We can read external data into a **DataFrame** using one of many **read_*** functions. We can also write from a **DataFrame** to a variety of formats using **to_*** functions. The most common of these are listed below:

Format type	Description	reader	writer
text	CSV	<code>read_csv</code>	<code>to_csv</code>
	Excel	<code>read_excel</code>	<code>to_excel</code>
text	JSON	<code>read_json</code>	<code>to_json</code>
binary	Feather	<code>read_feather</code>	<code>to_feather</code>
binary	SAS	<code>read_sas</code>	
SQL	SQL	<code>read_sql</code>	<code>to_sql</code>

We'll start by reading in the `mtcars` dataset stored as a CSV file

Note: You will need to have the data located in a file system that is exactly the same as mine for this to work. The `!ls`, which means list, line located below can be used to understand where you are located and navigate through your file system to get to where the directory called `data` is located.

```
[5]: !ls
```

```
00_python_primer.Rmd          08_genomics_project.ipynb
00_python_primer.ipynb        08_genomics_project_filledin.ipynb
01_python_tools_ds.Rmd        LICENSE
01_python_tools_ds.ipynb      README.md
02_python_pandas.ipynb        data
03_python_vis.ipynb           docs
04_python_stat.ipynb          graphs
05_practice_stats_filled.ipynb homeworks
05_stats_python_practice.ipynb live_coding
06_python_learning.ipynb      temporary.csv
07_python_appl.ipynb          workshop_documents
```

```
[6]: pd.read_csv('data/mtcars.csv')
```

```
[6]:
```

	make	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	\
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	
5	Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	
6	Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	
7	Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	
8	Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	
9	Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	
10	Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	
11	Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	
12	Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	
13	Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	
14	Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	
15	Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	
16	Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	
17	Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	
18	Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	
19	Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	
20	Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	
21	Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	
22	AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	
23	Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	

24	Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0
25	Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1
26	Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1
27	Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1
28	Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1
29	Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1
30	Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1
31	Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1

	gear	carb
0	4	4
1	4	4
2	4	1
3	3	1
4	3	2
5	3	1
6	3	4
7	4	2
8	4	2
9	4	4
10	4	4
11	3	3
12	3	3
13	3	3
14	3	4
15	3	4
16	3	4
17	4	1
18	4	2
19	4	1
20	3	1
21	3	2
22	3	2
23	3	4
24	3	2
25	4	1
26	5	2
27	5	2
28	5	4
29	5	6
30	5	8
31	4	2

This just prints out the data, but then it's lost. To use this data, we have to give it a name, so it's stored in Python's memory

```
[7]: mtcars = pd.read_csv('data/mtcars.csv')
```

One of the big differences between a spreadsheet program and a programming language from the data science perspective is that you have to load data into the programming language. It's not "just there" like Excel. This is a good thing, since it allows the common functionality of the programming language to work across multiple data sets, and also keeps the original data set pristine. Excel users can run into problems and [corrupt their data](#) if they are not careful.

If we wanted to write this data set back out into an Excel file, say, we could do

```
[8]: mtcars.to_excel('data/mtcars.xlsx')
```

You may get an error if you don't have the `openpyxl` package installed. You can easily install it from the Anaconda prompt using `conda install openpyxl` and following the prompts.

1.4 Exploring a data set

We would like to get some idea about this data set. There are a bunch of functions linked to the `DataFrame` object that help us in this. First we will use `head` to see the first 8 rows of this data set

```
[9]: mtcars.head(8)
```

```
[9]:
```

	make	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	\
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	
5	Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	
6	Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	
7	Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	


```
carb
```

0	4
1	4
2	1
3	1
4	2
5	1
6	4
7	2

This is our first look into this data. We notice a few things. Each column has a name, and each row has an *index*, starting at 0.

If you're interested in the last N rows, there is a corresponding `tail` function

Let's look at the data types of each of the columns

```
[10]: mtcars.dtypes
```

```
[10]: make      object
      mpg      float64
      cyl      int64
      disp     float64
      hp       int64
      drat     float64
      wt       float64
      qsec     float64
      vs       int64
      am       int64
      gear     int64
      carb     int64
      dtype: object
```

This tells us that some of the variables, like `mpg` and `disp`, are floating point (decimal) numbers, several are integers, and `make` is an “object”. The `dtypes` function borrows from `numpy`, where there isn’t really a type for character or categorical variables. So most often, when you see “object” in the output of `dtypes`, you think it’s a character or categorical variable.

We can also look at the data structure in a bit more detail.

```
[11]: mtcars.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32 entries, 0 to 31
Data columns (total 12 columns):
#   Column  Non-Null Count  Dtype
---  -
0   make    32 non-null      object
1   mpg     32 non-null      float64
2   cyl     32 non-null      int64
3   disp    32 non-null      float64
4   hp      32 non-null      int64
5   drat    32 non-null      float64
6   wt      32 non-null      float64
7   qsec    32 non-null      float64
8   vs      32 non-null      int64
9   am      32 non-null      int64
10  gear    32 non-null      int64
11  carb    32 non-null      int64
dtypes: float64(5), int64(6), object(1)
memory usage: 3.1+ KB
```

This tells us that this is indeed a `DataFrame`, with 12 columns, each with 32 valid observations. Each row has an index value ranging from 0 to 31. We also get the approximate size of this object in memory.

You can also quickly find the number of rows and columns of a data set by using `shape`, which is borrowed from `numpy`.

```
[12]: mtcars.shape
```

```
[12]: (32, 12)
```

More generally, we can get a summary of each variable using the `describe` function

```
[13]: mtcars.describe()
```

```
[13]:
```

	mpg	cyl	disp	hp	drat	wt \
count	32.000000	32.000000	32.000000	32.000000	32.000000	32.000000
mean	20.090625	6.187500	230.721875	146.687500	3.596563	3.217250
std	6.026948	1.785922	123.938694	68.562868	0.534679	0.978457
min	10.400000	4.000000	71.100000	52.000000	2.760000	1.513000
25%	15.425000	4.000000	120.825000	96.500000	3.080000	2.581250
50%	19.200000	6.000000	196.300000	123.000000	3.695000	3.325000
75%	22.800000	8.000000	326.000000	180.000000	3.920000	3.610000
max	33.900000	8.000000	472.000000	335.000000	4.930000	5.424000

	qsec	vs	am	gear	carb
count	32.000000	32.000000	32.000000	32.000000	32.0000
mean	17.848750	0.437500	0.406250	3.687500	2.8125
std	1.786943	0.504016	0.498991	0.737804	1.6152
min	14.500000	0.000000	0.000000	3.000000	1.0000
25%	16.892500	0.000000	0.000000	3.000000	2.0000
50%	17.710000	0.000000	0.000000	4.000000	2.0000
75%	18.900000	1.000000	1.000000	4.000000	4.0000
max	22.900000	1.000000	1.000000	5.000000	8.0000

These are usually the first steps in exploring the data.

1.5 Data structures and types

pandas has two main data types: **Series** and **DataFrame**. These are analogous to vectors and matrices, in that a **Series** is 1-dimensional while a **DataFrame** is 2-dimensional.

1.5.1 pandas.Series

The **Series** object holds data from a single input variable, and is required, much like numpy arrays, to be homogeneous in type. You can create **Series** objects from lists or numpy arrays quite easily

```
[14]: s = pd.Series([1,3,5,np.nan, 9, 13])
s
```

```
[14]: 0    1.0
      1    3.0
      2    5.0
      3   NaN
      4    9.0
      5   13.0
```

dtype: float64

```
[15]: s2 = pd.Series(np.arange(1,20))  
s2
```

```
[15]: 0      1  
      1      2  
      2      3  
      3      4  
      4      5  
      5      6  
      6      7  
      7      8  
      8      9  
      9     10  
     10     11  
     11     12  
     12     13  
     13     14  
     14     15  
     15     16  
     16     17  
     17     18  
     18     19  
dtype: int64
```

You can access elements of a **Series** much like a dict

```
[16]: s2[4]
```

```
[16]: 5
```

There is no requirement that the index of a **Series** has to be numeric. It can be any kind of scalar object

```
[17]: s3 = pd.Series(np.random.normal(0,1, (5,)), index = ['a','b','c','d','e'])  
s3
```

```
[17]: a      0.495129  
      b     -0.123514  
      c      0.779712  
      d     -1.216104  
      e      0.910280  
dtype: float64
```

```
[18]: s3['d']
```

```
[18]: -1.216103950445352
```

```
[19]: s3['a':'d']
```

```
[19]: a    0.495129  
      b   -0.123514  
      c    0.779712  
      d   -1.216104  
      dtype: float64
```

Well, slicing worked, but it gave us something different than expected. It gave us both the start **and** end of the slice, which is unlike what we've encountered so far!!

It turns out that in **pandas**, slicing by index actually does this. It is a discrepancy from **numpy** and Python in general that we have to be careful about.

You can extract the actual values into a numpy array

```
[20]: s3.to_numpy()
```

```
[20]: array([ 0.49512917, -0.1235141 ,  0.77971247, -1.21610395,  0.91027988])
```

In fact, you'll see that much of **pandas**' structures are build on top of **numpy** arrays. This is a good thing, since you can take advantage of the powerful numpy functions that are built for fast, efficient scientific computing.

Making the point about slicing again,

```
[21]: s3.to_numpy()[0:3]
```

```
[21]: array([ 0.49512917, -0.1235141 ,  0.77971247])
```

This is different from index-based slicing done earlier.

1.5.2 pandas.DataFrame

The **DataFrame** object holds a rectangular data set. Each column of a **DataFrame** is a **Series** object. This means that each column of a **DataFrame** must be comprised of data of the same type, but different columns can hold data of different types. This structure is extremely useful in practical data science. The invention of this structure was, in my opinion, transformative in making Python an effective data science tool.

Creating a DataFrame The **DataFrame** can be created by importing data, as we saw in the previous section. It can also be created by a few methods within Python.

First, it can be created from a 2-dimensional **numpy** array.

```
[22]: rng = np.random.RandomState(25)  
      d1 = pd.DataFrame(rng.normal(0,1, (4,5)))  
      d1
```



```
[22]:
```

	0	1	2	3	4
0	0.228273	1.026890	-0.839585	-0.591182	-0.956888
1	-0.222326	-0.619915	1.837905	-2.053231	0.868583
2	-0.920734	-0.232312	2.152957	-1.334661	0.076380
3	-1.246089	1.202272	-1.049942	1.056610	-0.419678

You will notice that it creates default column names, that are merely the column number, starting from 0. We can also create the column names and row index (similar to the `Series` index we saw earlier) directly during creation.

```
[23]: d2 = pd.DataFrame(rng.normal(0,1, (4, 5)),
                        columns = ['A', 'B', 'C', 'D', 'E'],
                        index = ['a', 'b', 'c', 'd'])
d2
```

```
[23]:
```

	A	B	C	D	E
a	2.294842	-2.594487	2.822756	0.680889	-1.577693
b	-1.976254	0.533340	-0.290870	-0.513520	1.982626
c	0.226001	-1.839905	1.607671	0.388292	0.399732
d	0.405477	0.217002	-0.633439	0.246622	-1.939546

We could also create a `DataFrame` from a list of lists, as long as things line up, just as we showed for `numpy` arrays. However, to me, other ways, including the `dict` method below, make more sense.

We can change the column names (which can be extracted and replaced with the `columns` attribute) and the index values (using the `index` attribute).

```
[24]: d2.columns
```

```
[24]: Index(['A', 'B', 'C', 'D', 'E'], dtype='object')
```

```
[25]: d2.columns = pd.Index(['V'+str(i) for i in range(1,6)]) # Index creates the
↳right objects for both column names and row names, which can be extracted
↳and changed with the `index` attribute
d2
```

```
[25]:
```

	V1	V2	V3	V4	V5
a	2.294842	-2.594487	2.822756	0.680889	-1.577693
b	-1.976254	0.533340	-0.290870	-0.513520	1.982626
c	0.226001	-1.839905	1.607671	0.388292	0.399732
d	0.405477	0.217002	-0.633439	0.246622	-1.939546

Exercise: Can you explain what I did in the list comprehension above? The key points are understanding `str` and how I constructed the `range`.

```
[26]: d2.index = ['o1', 'o2', 'o3', 'o4']
d2
```

```
[26]:
```

	V1	V2	V3	V4	V5
o1	2.294842	-2.594487	2.822756	0.680889	-1.577693
o2	-1.976254	0.533340	-0.290870	-0.513520	1.982626
o3	0.226001	-1.839905	1.607671	0.388292	0.399732
o4	0.405477	0.217002	-0.633439	0.246622	-1.939546

You can also extract data from a homogeneous `DataFrame` to a `numpy` array

```
[27]: d1.to_numpy()
```

```
[27]: array([[ 0.22827309,  1.0268903 , -0.83958485, -0.59118152, -0.9568883 ],
             [-0.22232569, -0.61991511,  1.83790458, -2.05323076,  0.86858305],
             [-0.92073444, -0.23231186,  2.1529569 , -1.33466147,  0.07637965],
             [-1.24608928,  1.20227231, -1.04994158,  1.05661011, -0.41967767]])
```

It turns out that you can use `to_numpy` for a non-homogeneous `DataFrame` as well. `numpy` just makes it homogeneous by assigning each column the data type `object`. This also limits what you can do in `numpy` with the array and may require changing data types using the `astype` function. There is some more detail about the `object` data type in the Python Tools for Data Science ([notebook](#), [PDF](#)) document.

The other easy way to create a `DataFrame` is from a `dict` object, where each component object is either a list or a `numpy` array, and is homogeneous in type. One exception is if a component is of size 1; then it is repeated to meet the needs of the `DataFrame`'s dimensions

```
[28]: df = pd.DataFrame({
    'A': 3.,
    'B': rng.random_sample(5),
    'C': pd.Timestamp('20200512'),
    'D': np.array([6] * 5),
    'E': pd.Categorical(['yes', 'no', 'no', 'yes', 'no']),
    'F': 'NIH'})
df
```

```
[28]:
```

	A	B	C	D	E	F
0	3.0	0.958092	2020-05-12	6	yes	NIH
1	3.0	0.883201	2020-05-12	6	no	NIH
2	3.0	0.295432	2020-05-12	6	no	NIH
3	3.0	0.512376	2020-05-12	6	yes	NIH
4	3.0	0.088702	2020-05-12	6	no	NIH

```
[29]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5 entries, 0 to 4
Data columns (total 6 columns):
#   Column  Non-Null Count  Dtype
---  -
0    A      5 non-null         float64
```

```

1   B      5 non-null    float64
2   C      5 non-null    datetime64[ns]
3   D      5 non-null    int64
4   E      5 non-null    category
5   F      5 non-null    object
dtypes: category(1), datetime64[ns](1), float64(2), int64(1), object(1)
memory usage: 429.0+ bytes

```

We note that C is a date object, E is a category object, and F is a text/string object. pandas has excellent time series capabilities (having origins in FinTech), and the `TimeStamp` function creates datetime objects which can be queried and manipulated in Python. We'll describe category data in the next section.

You can also create a `DataFrame` where each column is composed of composite objects, like lists and dicts, as well. This might have limited value in some settings, but may be useful in others. In particular, this allows capabilities like the *list-column construct in R tibbles*. For example,

```
[30]: pd.DataFrame({'list' : [[1,2],[3,4],[5,6]],
                    'tuple' : [('a','b'), ('c','d'), ('e','f')],
                    'set'   : [{ 'A', 'B', 'C' }, { 'D', 'E' }, { 'F' }],
                    'dicts' : [{ 'A': [1,2,3]}, { 'B': [5,6,8]}, { 'C': [3,9]}]})

```

```
[30]:      list  tuple      set      dicts
0  [1, 2]  (a, b)  {A, C, B}  {'A': [1, 2, 3]}
1  [3, 4]  (c, d)   {E, D}  {'B': [5, 6, 8]}
2  [5, 6]  (e, f)    {F}   {'C': [3, 9]}

```

Working with a DataFrame You can extract particular columns of a `DataFrame` by name

```
[31]: df['E']

```

```
[31]: 0    yes
      1    no
      2    no
      3    yes
      4    no
      Name: E, dtype: category
      Categories (2, object): [no, yes]

```

```
[32]: df['B']

```

```
[32]: 0    0.958092
      1    0.883201
      2    0.295432
      3    0.512376
      4    0.088702
      Name: B, dtype: float64

```

There is also a shortcut for accessing single columns, using Python's dot (`.`) notation.

```
[33]: df.B
```

```
[33]: 0    0.958092
      1    0.883201
      2    0.295432
      3    0.512376
      4    0.088702
      Name: B, dtype: float64
```

This notation can be more convenient if we need to perform operations on a single column. If we want to extract multiple columns, this notation will not work. Also, if we want to create new columns or replace existing columns, we need to use the array notation with the column name in quotes.

Let's look at slicing a `DataFrame`

Extracting rows and columns There are two extractor functions in `pandas`:

- `loc` extracts by label (index label, column label, slice of labels, etc.
- `iloc` extracts by index (integers, slice objects, etc.

```
[34]: df2 = pd.DataFrame(rng.randint(0,10, (5,4)),
                        index = ['a','b','c','d','e'],
                        columns = ['one','two','three','four'])
df2
```

```
[34]:   one  two  three  four
a     5   3     2     8
b     9   3     0     5
c     8   4     3     3
d     5   2     7     1
e     6   7     8     7
```

First, let's see what naively slicing this `DataFrame` does.

```
[35]: df2['one']
```

```
[35]: a     5
      b     9
      c     8
      d     5
      e     6
      Name: one, dtype: int64
```

Ok, that works. It grabs one column from the dataset. How about the dot notation?

```
[36]: df2.one
```

```
[36]: a    5
      b    9
      c    8
      d    5
      e    6
      Name: one, dtype: int64
```

Let's see what this produces.

```
[37]: type(df2.one)
```

```
[37]: pandas.core.series.Series
```

So this is a series, so we can potentially do slicing of this series.

```
[38]: df2.one['b']
```

```
[38]: 9
```

```
[39]: df2.one['b':'d']
```

```
[39]: b    9
      c    8
      d    5
      Name: one, dtype: int64
```

```
[40]: df2.one[:3]
```

```
[40]: a    5
      b    9
      c    8
      Name: one, dtype: int64
```

Ok, so we have all the **Series** slicing available. The problem here is in semantics, in that we are grabbing one column and then slicing the rows. That doesn't quite work with our sense that a **DataFrame** is a rectangle with rows and columns, and we tend to think of rows, then columns.

Let's see if we can do column slicing with this.

```
[41]: df2[:, 'two']
```

```
[41]:   one  two  three  four
a     5    3     2     8
b     9    3     0     5
c     8    4     3     3
d     5    2     7     1
e     6    7     8     7
```

That's not what we want, of course. It's giving back the entire data frame. We'll come back to this.

```
[42]: df2[['one', 'three']]
```

```
[42]:
```

	one	three
a	5	2
b	9	0
c	8	3
d	5	7
e	6	8

That works correctly though. We can give a list of column names. Ok.

How about row slices?

```
[43]: #df2['a'] # Doesn't work
df2['a':'c']
```

```
[43]:
```

	one	two	three	four
a	5	3	2	8
b	9	3	0	5
c	8	4	3	3

Ok, that works. It slices rows, but includes the largest index, like a **Series** but unlike **numpy** arrays.

```
[44]: df2[0:2]
```

```
[44]:
```

	one	two	three	four
a	5	3	2	8
b	9	3	0	5

Slices by location work too, but use the **numpy** slicing rules.

This entire extraction method becomes confusing. Let's simplify things for this, and then move on to more consistent ways to extract elements of a **DataFrame**. Let's agree on two things. If we're going the direct extraction route,

1. We will extract single columns of a **DataFrame** with `[]` or `.`, i.e., `df2['one']` or `df2.one`
2. We will extract slices of rows of a **DataFrame** using location only, i.e., `df2[:3]`.

For everything else, we'll use two functions, `loc` and `iloc`.

- `loc` extracts elements like a matrix, using index and columns
- `iloc` extracts elements like a matrix, using location

```
[45]: df2.loc[:, 'one': 'three']
```

```
[45]:
```

	one	two	three
a	5	3	2
b	9	3	0

c	8	4	3
d	5	2	7
e	6	7	8

```
[46]: df2.loc['a':'d',:]
```

```
[46]:
```

	one	two	three	four
a	5	3	2	8
b	9	3	0	5
c	8	4	3	3
d	5	2	7	1

```
[47]: df2.loc['b', 'three']
```

```
[47]: 0
```

So loc works just like a matrix, but with **pandas** slicing rules (include largest index)

```
[48]: df2.iloc[:,1:4]
```

```
[48]:
```

	two	three	four
a	3	2	8
b	3	0	5
c	4	3	3
d	2	7	1
e	7	8	7

```
[49]: df2.iloc[1:3,:]
```

```
[49]:
```

	one	two	three	four
b	9	3	0	5
c	8	4	3	3

```
[50]: df2.iloc[1:3, 1:4]
```

```
[50]:
```

	two	three	four
b	3	0	5
c	4	3	3

iloc slices like a matrix, but uses **numpy** slicing conventions (does **not** include highest index)

If we want to extract a single element from a dataset, there are two functions available, **iat** and **at**, with behavior corresponding to **iloc** and **loc**, respectively.

```
[51]: df2.iat[2,3]
```

```
[51]: 3
```

```
[52]: df2.at['b', 'three']
```

```
[52]: 0
```

Boolean selection We can also use tests to extract data from a `DataFrame`. For example, we can extract only rows where column labeled `one` is greater than 3.

```
[53]: df2[df2.one > 3]
```

```
[53]:
```

	one	two	three	four
a	5	3	2	8
b	9	3	0	5
c	8	4	3	3
d	5	2	7	1
e	6	7	8	7

We can also do composite tests. Here we ask for rows where `one` is greater than 3 and `three` is less than 9

```
[54]: df2[(df2.one > 3) & (df2.three < 9)]
```

```
[54]:
```

	one	two	three	four
a	5	3	2	8
b	9	3	0	5
c	8	4	3	3
d	5	2	7	1
e	6	7	8	7

query `DataFrame`'s have a `query` method allowing selection using a Python expression

```
[55]: n = 10
df = pd.DataFrame(np.random.rand(n, 3), columns = list('abc'))
df
```

```
[55]:
```

	a	b	c
0	0.815508	0.028886	0.785901
1	0.300541	0.338317	0.641448
2	0.505391	0.112547	0.185908
3	0.400225	0.931705	0.094609
4	0.981477	0.008321	0.880875
5	0.876385	0.144667	0.768066
6	0.000640	0.269403	0.649907
7	0.250573	0.479758	0.654926
8	0.358420	0.624067	0.374675
9	0.686182	0.661075	0.398815

```
[56]: df[(df.a < df.b) & (df.b < df.c)]
```



```
[56]:
```

	a	b	c
1	0.300541	0.338317	0.641448
6	0.000640	0.269403	0.649907
7	0.250573	0.479758	0.654926

We can equivalently write this query as

```
[57]: df.query('(a < b) & (b < c)')
```

```
[57]:
```

	a	b	c
1	0.300541	0.338317	0.641448
6	0.000640	0.269403	0.649907
7	0.250573	0.479758	0.654926

Replacing values in a DataFrame We can replace values within a DataFrame either by position or using a query.

```
[58]: df2
```

```
[58]:
```

	one	two	three	four
a	5	3	2	8
b	9	3	0	5
c	8	4	3	3
d	5	2	7	1
e	6	7	8	7

```
[59]: df2['one'] = [2,5,2,5,2]
df2
```

```
[59]:
```

	one	two	three	four
a	2	3	2	8
b	5	3	0	5
c	2	4	3	3
d	5	2	7	1
e	2	7	8	7

```
[60]: df2.iat[2,3] = -9 # missing value
df2
```

```
[60]:
```

	one	two	three	four
a	2	3	2	8
b	5	3	0	5
c	2	4	3	-9
d	5	2	7	1
e	2	7	8	7

Let's now replace values using `replace` which is more flexible.

```
[61]: df2.replace(0, -9) # replace 0 with -9
```

```
[61]:
```

	one	two	three	four
a	2	3	2	8
b	5	3	-9	5
c	2	4	3	-9
d	5	2	7	1
e	2	7	8	7

```
[62]: df2.replace({2: 2.5, 8: 6.5}) # multiple replacements
```

```
[62]:
```

	one	two	three	four
a	2.5	3.0	2.5	6.5
b	5.0	3.0	0.0	5.0
c	2.5	4.0	3.0	-9.0
d	5.0	2.5	7.0	1.0
e	2.5	7.0	6.5	7.0

```
[63]: df2.replace({'one': {5: 500}, 'three': {0: -9, 8: 800}})  
# different replacements in different columns
```

```
[63]:
```

	one	two	three	four
a	2	3	2	8
b	500	3	-9	5
c	2	4	3	-9
d	500	2	7	1
e	2	7	800	7

See more examples in the [documentation](#)

1.5.3 Categorical data

pandas provides a `Categorical` function and a `category` object type to Python. This type is analogous to the `factor` data type in R. It is meant to address categorical or discrete variables, where we need to use them in analyses. Categorical variables typically take on a small number of unique values, like gender, blood type, country of origin, race, etc.

You can create categorical `Series` in a couple of ways:

```
[64]: s = pd.Series(['a', 'b', 'c'], dtype='category')
```

```
[65]: df = pd.DataFrame({  
    'A': 3.,  
    'B': rng.random_sample(5),  
    'C': pd.Timestamp('20200512'),  
    'D': np.array([6] * 5),  
    'E': pd.Categorical(['yes', 'no', 'no', 'yes', 'no']),  
    'F': 'NIH'})
```

```
df['F'].astype('category')
```

```
[65]: 0    NIH
      1    NIH
      2    NIH
      3    NIH
      4    NIH
      Name: F, dtype: category
      Categories (1, object): [NIH]
```

You can also create DataFrame's where each column is categorical

```
[66]: df = pd.DataFrame({'A': list('abcd'), 'B': list('bdca')})
      df_cat = df.astype('category')
      df_cat.dtypes
```

```
[66]: A    category
      B    category
      dtype: object
```

You can explore categorical data in a variety of ways

```
[67]: df_cat['A'].describe()
```

```
[67]: count      4
      unique      4
      top         d
      freq        1
      Name: A, dtype: object
```

```
[68]: df['A'].value_counts()
```

```
[68]: a    1
      b    1
      c    1
      d    1
      Name: A, dtype: int64
```

One issue with categories is that, if a particular level of a category is not seen before, it can create an error. So you can pre-specify the categories you expect

```
[69]: df_cat['B'] = pd.Categorical(list('aabb'), categories = ['a','b','c','d'])
      df_cat['B'].value_counts()
```

```
[69]: b    2
      a    2
      d    0
      c    0
```

Name: B, dtype: int64

Re-organizing categories In categorical data, there is often the concept of a “first” or “reference” category, and an ordering of categories. This tends to be important in both visualization as well as in regression modeling. Both aspects of a category can be addressed using the `reorder_categories` function.

In our earlier example, we can see that the A variable has 4 categories, with the “first” category being “a”.

```
[70]: df_cat.A
```

```
[70]: 0    a
      1    b
      2    c
      3    d
      Name: A, dtype: category
      Categories (4, object): [a, b, c, d]
```

Suppose we want to change this ordering to the reverse ordering, where “d” is the “first” category, and then it goes in reverse order.

```
[71]: df_cat['A'] = df_cat.A.cat.reorder_categories(['d','c','b','a'])
      df_cat.A
```

```
[71]: 0    a
      1    b
      2    c
      3    d
      Name: A, dtype: category
      Categories (4, object): [d, c, b, a]
```

1.5.4 Missing data

Both `numpy` and `pandas` allow for missing values, which are a reality in data science. The missing values are coded as `np.nan`. Let’s create some data and force some missing values

```
[72]: df = pd.DataFrame(np.random.randn(5, 3), index = ['a','c','e', 'f','g'],
      ↪ columns = ['one','two','three']) # pre-specify index and column names
      df['four'] = 20 # add a column named "four", which will all be 20
      df['five'] = df['one'] > 0
      df
```

```
[72]:      one      two      three  four  five
a -0.461526  0.100114  1.203480    20  False
c  0.895374  0.411451 -0.090356    20   True
e  0.411101  1.047826 -1.688879    20   True
f  1.083412 -1.034765  0.002112    20   True
```

```
g 1.974888 0.255597 -0.949174 20 True
```

```
[73]: df2 = df.reindex(['a','b','c','d','e','f','g'])
df2.style.applymap(lambda x: 'background-color:yellow', subset = pd.
↳ IndexSlice[['b','d'],:])
```

```
[73]: <pandas.io.formats.style.Styler at 0x7feba82e1310>
```

The code above is creating new blank rows based on the new index values, some of which are present in the existing data and some of which are missing.

We can create *masks* of the data indicating where missing values reside in a data set.

```
[74]: df2.isna()
```

```
[74]:
```

	one	two	three	four	five
a	False	False	False	False	False
b	True	True	True	True	True
c	False	False	False	False	False
d	True	True	True	True	True
e	False	False	False	False	False
f	False	False	False	False	False
g	False	False	False	False	False

```
[75]: df2['one'].notna()
```

```
[75]: a      True
b     False
c      True
d     False
e      True
f      True
g      True
Name: one, dtype: bool
```

We can obtain complete data by dropping any row that has any missing value. This is called *complete case analysis*, and you should be very careful using it. It is *only* valid if we believe that the missingness is missing at random, and not related to some characteristic of the data or the data gathering process.

```
[76]: df2.dropna(how='any')
```

```
[76]:
```

	one	two	three	four	five
a	-0.461526	0.100114	1.203480	20.0	False
c	0.895374	0.411451	-0.090356	20.0	True
e	0.411101	1.047826	-1.688879	20.0	True
f	1.083412	-1.034765	0.002112	20.0	True
g	1.974888	0.255597	-0.949174	20.0	True

You can also fill in, or *impute*, missing values. This can be done using a single value..

```
[77]: out1 = df2.fillna(value = 5)

out1.style.applymap(lambda x: 'background-color:yellow', subset = pd.
↳IndexSlice[['b','d'],:])
```

```
[77]: <pandas.io.formats.style.Styler at 0x7feba8e3b290>
```

or a computed value like a column mean

```
[78]: df3 = df2.copy()
df3 = df3.select_dtypes(exclude=[object]) # remove non-numeric columns
out2 = df3.fillna(df3.mean()) # df3.mean() computes column-wise means

out2.style.applymap(lambda x: 'background-color:yellow', subset = pd.
↳IndexSlice[['b','d'],:])
```

```
[78]: <pandas.io.formats.style.Styler at 0x7feba82dd110>
```

You can also impute based on the principle of *last value carried forward* which is common in time series. This means that the missing value is imputed with the previous recorded value.

```
[79]: out3 = df2.fillna(method = 'ffill') # Fill forward

out3.style.applymap(lambda x: 'background-color:yellow', subset = pd.
↳IndexSlice[['b','d'],:])
```

```
[79]: <pandas.io.formats.style.Styler at 0x7feba9382390>
```

```
[80]: out4 = df2.fillna(method = 'bfill') # Fill backward

out4.style.applymap(lambda x: 'background-color:yellow', subset = pd.
↳IndexSlice[['b','d'],:])
```

```
[80]: <pandas.io.formats.style.Styler at 0x7feba9398e10>
```

1.6 Data transformation

1.6.1 Arithmetic operations

If you have a Series or DataFrame that is all numeric, you can add or multiply single numbers to all the elements together.

```
[81]: A = pd.DataFrame(np.random.randn(4,5))
print(A)
```

```
      0      1      2      3      4
0  0.670003  0.252635  0.440303  1.015211  2.682897
```

```

1 -0.255074 -1.039442  0.586182 -0.160948 -1.552889
2 -0.609655  0.540724  1.277570 -0.583918 -0.260800
3  2.814046  0.303757  0.170083  0.090630  0.568430

```

```
[82]: print(A + 6)
```

```

      0      1      2      3      4
0  6.670003  6.252635  6.440303  7.015211  8.682897
1  5.744926  4.960558  6.586182  5.839052  4.447111
2  5.390345  6.540724  7.277570  5.416082  5.739200
3  8.814046  6.303757  6.170083  6.090630  6.568430

```

```
[83]: print(A * -10)
```

```

      0      1      2      3      4
0 -6.700033 -2.526347 -4.403026 -10.152107 -26.828973
1  2.550739 10.394422 -5.861815  1.609480 15.528886
2  6.096548 -5.407241 -12.775704  5.839178  2.608003
3 -28.140456 -3.037571 -1.700833 -0.906296 -5.684301

```

If you have two compatible (same dimension) numeric `DataFrames`, you can add, subtract, multiply and divide elementwise

```
[84]: B = pd.DataFrame(np.random.randn(4,5) + 4)
      print(A + B)
```

```

      0      1      2      3      4
0  4.168630  3.687596  3.975981  5.392280  7.763671
1  3.658142  1.958768  5.279375  5.128594  0.011606
2  1.991331  5.950211  8.031543  2.998088  5.140077
3  7.058083  5.038382  2.142168  3.489493  5.321225

```

```
[85]: print(A * B)
```

```

      0      1      2      3      4
0  2.344091  0.867790  1.556768  4.443648 13.631194
1 -0.998160 -3.116466  2.751063 -0.851341 -2.429485
2 -1.585703  2.925040  8.628675 -2.091597 -1.408551
3 11.942915  1.438176  0.335419  0.308038  2.701632

```

If you have a `Series` with the same number of elements as the number of columns of a `DataFrame`, you can do arithmetic operations, with each element of the `Series` acting upon each column of the `DataFrame`

```
[86]: c = pd.Series([1,2,3,4,5])
      print(A + c)
```

```

      0      1      2      3      4
0  1.670003  2.252635  3.440303  5.015211  7.682897
1  0.744926  0.960558  3.586182  3.839052  3.447111

```

```

2  0.390345  2.540724  4.277570  3.416082  4.739200
3  3.814046  2.303757  3.170083  4.090630  5.568430

```

```
[87]: print(A * c)
```

```

      0      1      2      3      4
0  0.670003  0.505269  1.320908  4.060843  13.414487
1 -0.255074 -2.078884  1.758545 -0.643792 -7.764443
2 -0.609655  1.081448  3.832711 -2.335671 -1.304001
3  2.814046  0.607514  0.510250  0.362519  2.842150

```

This idea can be used to standardize a dataset, i.e. make each column have mean 0 and standard deviation 1.

```
[88]: means = A.mean(axis=0)
stds = A.std(axis = 0)

(A - means)/stds
```

```

[88]:      0      1      2      3      4
0  0.009870  0.333779 -0.377639  1.367184  1.306800
1 -0.591909 -1.476625 -0.068550 -0.371284 -1.075535
2 -0.822570  0.737437  1.396369 -0.996471 -0.348825
3  1.404608  0.405409 -0.950181  0.000571  0.117559

```

1.6.2 Concatenation of data sets

Let's create some example data sets

```
[89]: df1 = pd.DataFrame({'A': ['a'+str(i) for i in range(4)],
    'B': ['b'+str(i) for i in range(4)],
    'C': ['c'+str(i) for i in range(4)],
    'D': ['d'+str(i) for i in range(4)]})

df2 = pd.DataFrame({'A': ['a'+str(i) for i in range(4,8)],
    'B': ['b'+str(i) for i in range(4,8)],
    'C': ['c'+str(i) for i in range(4,8)],
    'D': ['d'+str(i) for i in range(4,8)]})

df3 = pd.DataFrame({'A': ['a'+str(i) for i in range(8,12)],
    'B': ['b'+str(i) for i in range(8,12)],
    'C': ['c'+str(i) for i in range(8,12)],
    'D': ['d'+str(i) for i in range(8,12)]})
```

We can concatenate these DataFrame objects by row

```
[90]: row_concatenate = pd.concat([df1, df2, df3])
print(row_concatenate)
```

```

      A      B      C      D

```


0	a0	b0	c0	d0
1	a1	b1	c1	d1
2	a2	b2	c2	d2
3	a3	b3	c3	d3
0	a4	b4	c4	d4
1	a5	b5	c5	d5
2	a6	b6	c6	d6
3	a7	b7	c7	d7
0	a8	b8	c8	d8
1	a9	b9	c9	d9
2	a10	b10	c10	d10
3	a11	b11	c11	d11

This stacks the dataframes together. They are literally stacked, as is evidenced by the index values being repeated.

This same exercise can be done by the **append** function

```
[91]: df1.append(df2).append(df3)
```

```
[91]:
```

	A	B	C	D
0	a0	b0	c0	d0
1	a1	b1	c1	d1
2	a2	b2	c2	d2
3	a3	b3	c3	d3
0	a4	b4	c4	d4
1	a5	b5	c5	d5
2	a6	b6	c6	d6
3	a7	b7	c7	d7
0	a8	b8	c8	d8
1	a9	b9	c9	d9
2	a10	b10	c10	d10
3	a11	b11	c11	d11

Suppose we want to append a new row to **df1**. Lets create a new row.

```
[92]: new_row = pd.Series(['n1', 'n2', 'n3', 'n4'])
      pd.concat([df1, new_row])
```

```
[92]:
```

	A	B	C	D	0
0	a0	b0	c0	d0	NaN
1	a1	b1	c1	d1	NaN
2	a2	b2	c2	d2	NaN
3	a3	b3	c3	d3	NaN
0	NaN	NaN	NaN	NaN	n1
1	NaN	NaN	NaN	NaN	n2
2	NaN	NaN	NaN	NaN	n3
3	NaN	NaN	NaN	NaN	n4

That's a lot of missing values. The issue is that the we don't have column names in the `new_row`, and the indices are the same, so pandas tries to append it my making a new column. The solution is to make it a `DataFrame`.

```
[93]: new_row = pd.DataFrame(['n1', 'n2', 'n3', 'n4'], columns = ['A', 'B', 'C', 'D'])
      print(new_row)
```

```
      A  B  C  D
0  n1  n2  n3  n4
```

```
[94]: pd.concat([df1, new_row])
```

```
[94]:      A  B  C  D
0  a0  b0  c0  d0
1  a1  b1  c1  d1
2  a2  b2  c2  d2
3  a3  b3  c3  d3
0  n1  n2  n3  n4
```

or

```
[95]: df1.append(new_row)
```

```
[95]:      A  B  C  D
0  a0  b0  c0  d0
1  a1  b1  c1  d1
2  a2  b2  c2  d2
3  a3  b3  c3  d3
0  n1  n2  n3  n4
```

Adding columns

```
[96]: pd.concat([df1, df2, df3], axis = 1)
```

```
[96]:      A  B  C  D  A  B  C  D  A  B  C  D
0  a0  b0  c0  d0  a4  b4  c4  d4  a8  b8  c8  d8
1  a1  b1  c1  d1  a5  b5  c5  d5  a9  b9  c9  d9
2  a2  b2  c2  d2  a6  b6  c6  d6  a10 b10 c10 d10
3  a3  b3  c3  d3  a7  b7  c7  d7  a11 b11 c11 d11
```

The option `axis=1` ensures that concatenation happens by columns. The default value `axis = 0` concatenates by rows.

Let's play a little game. Let's change the column names of `df2` and `df3` so they are not the same as `df1`.

```
[97]: df2.columns = ['E', 'F', 'G', 'H']
      df3.columns = ['A', 'D', 'F', 'H']
      pd.concat([df1, df2, df3])
```

```
[97]:
```

	A	B	C	D	E	F	G	H
0	a0	b0	c0	d0	NaN	NaN	NaN	NaN
1	a1	b1	c1	d1	NaN	NaN	NaN	NaN
2	a2	b2	c2	d2	NaN	NaN	NaN	NaN
3	a3	b3	c3	d3	NaN	NaN	NaN	NaN
0	NaN	NaN	NaN	NaN	a4	b4	c4	d4
1	NaN	NaN	NaN	NaN	a5	b5	c5	d5
2	NaN	NaN	NaN	NaN	a6	b6	c6	d6
3	NaN	NaN	NaN	NaN	a7	b7	c7	d7
0	a8	NaN	NaN	b8	NaN	c8	NaN	d8
1	a9	NaN	NaN	b9	NaN	c9	NaN	d9
2	a10	NaN	NaN	b10	NaN	c10	NaN	d10
3	a11	NaN	NaN	b11	NaN	c11	NaN	d11

Now pandas ensures that all column names are represented in the new data frame, but with missing values where the row indices and column indices are mismatched. Some of this can be avoided by only joining on common columns. This is done using the `join` option in `concat`. The default value is 'outer', which is what you see. above

```
[98]: pd.concat([df1, df3], join = 'inner')
```

```
[98]:
```

	A	D
0	a0	d0
1	a1	d1
2	a2	d2
3	a3	d3
0	a8	b8
1	a9	b9
2	a10	b10
3	a11	b11

You can do the same thing when joining by rows, using `axis = 0` and `join="inner"` to only join on rows with matching indices. Reminder that the indices are just labels and happen to be the row numbers by default.

1.6.3 Merging data sets

For this section we'll use a set of data from a survey, also used by Daniel Chen in "Pandas for Everyone"

```
[99]: person = pd.read_csv('data/survey_person.csv')
      site = pd.read_csv('data/survey_site.csv')
      survey = pd.read_csv('data/survey_survey.csv')
      visited = pd.read_csv('data/survey_visited.csv')
```

```
[100]: print(person)
```

	ident	personal	family
0	dyer	William	Dyer

1	pb	Frank	Pabodie
2	lake	Anderson	Lake
3	roe	Valentina	Roerich
4	danforth	Frank	Danforth

```
[101]: print(site)
```

	name	lat	long
0	DR-1	-49.85	-128.57
1	DR-3	-47.15	-126.72
2	MSK-4	-48.87	-123.40

```
[129]: print(survey)
```

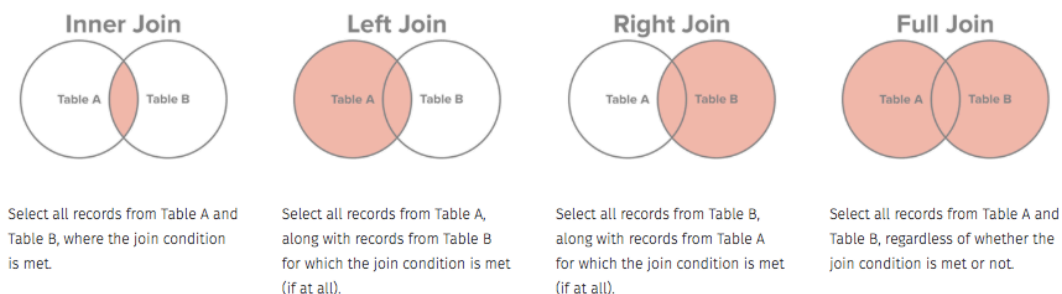
	taken	person	quant	reading
0	619	dyer	rad	9.82
1	619	dyer	sal	0.13
2	622	dyer	rad	7.80
3	622	dyer	sal	0.09
4	734	pb	rad	8.41
5	734	lake	sal	0.05
6	734	pb	temp	-21.50
7	735	pb	rad	7.22
8	735	NaN	sal	0.06
9	735	NaN	temp	-26.00
10	751	pb	rad	4.35
11	751	pb	temp	-18.50
12	751	lake	sal	0.10
13	752	lake	rad	2.19
14	752	lake	sal	0.09
15	752	lake	temp	-16.00
16	752	roe	sal	41.60
17	837	lake	rad	1.46
18	837	lake	sal	0.21
19	837	roe	sal	22.50
20	844	roe	rad	11.25

```
[102]: print(visited)
```

	ident	site	dated
0	619	DR-1	1927-02-08
1	622	DR-1	1927-02-10
2	734	DR-3	1939-01-07
3	735	DR-3	1930-01-12
4	751	DR-3	1930-02-26
5	752	DR-3	NaN
6	837	MSK-4	1932-01-14
7	844	DR-1	1932-03-22

There are basically four kinds of joins:

pandas	R	SQL	Description
left	left_join	left outer	keep all rows on left
right	right_join	right outer	keep all rows on right
outer	outer_join	full outer	keep all rows from both
inner	inner_join	inner	keep only rows with common keys



The terms `left` and `right` refer to which data set you call first and second respectively.

We start with an left join

```
[103]: s2v_merge = survey.merge(visited, left_on = 'taken',right_on = 'ident', how =
      ↪ 'left')
```

```
[104]: print(s2v_merge)
```

```
   taken person quant  reading  ident  site      dated
0    619   dyer   rad    9.82    619  DR-1  1927-02-08
1    619   dyer   sal    0.13    619  DR-1  1927-02-08
2    622   dyer   rad    7.80    622  DR-1  1927-02-10
3    622   dyer   sal    0.09    622  DR-1  1927-02-10
4    734    pb   rad    8.41    734  DR-3  1939-01-07
5    734   lake   sal    0.05    734  DR-3  1939-01-07
6    734    pb  temp   -21.50    734  DR-3  1939-01-07
7    735    pb   rad    7.22    735  DR-3  1930-01-12
8    735   NaN   sal    0.06    735  DR-3  1930-01-12
9    735   NaN  temp   -26.00    735  DR-3  1930-01-12
10   751    pb   rad    4.35    751  DR-3  1930-02-26
11   751    pb  temp   -18.50    751  DR-3  1930-02-26
12   751   lake   sal    0.10    751  DR-3  1930-02-26
13   752   lake   rad    2.19    752  DR-3          NaN
14   752   lake   sal    0.09    752  DR-3          NaN
15   752   lake  temp   -16.00    752  DR-3          NaN
16   752    roe   sal    41.60    752  DR-3          NaN
17   837   lake   rad    1.46    837  MSK-4  1932-01-14
18   837   lake   sal    0.21    837  MSK-4  1932-01-14
```

19	837	roe	sal	22.50	837	MSK-4	1932-01-14
20	844	roe	rad	11.25	844	DR-1	1932-03-22

Here, the left dataset is `survey` and the right one is `visited`. Since we're doing a left join, we keep all the rows from `survey` and add columns from `visited`, matching on the common key, called "taken" in one dataset and "ident" in the other. Note that the rows of `visited` are repeated as needed to line up with all the rows with common "taken" values.

We can now add location information, where the common key is the site code

```
[105]: s2v2loc_merge = s2v_merge.merge(site, how = 'left', left_on = 'site', right_on='
      ↪= 'name')
      print(s2v2loc_merge)
```

	taken	person	quant	reading	ident	site	dated	name	lat	\
0	619	dyer	rad	9.82	619	DR-1	1927-02-08	DR-1	-49.85	
1	619	dyer	sal	0.13	619	DR-1	1927-02-08	DR-1	-49.85	
2	622	dyer	rad	7.80	622	DR-1	1927-02-10	DR-1	-49.85	
3	622	dyer	sal	0.09	622	DR-1	1927-02-10	DR-1	-49.85	
4	734	pb	rad	8.41	734	DR-3	1939-01-07	DR-3	-47.15	
5	734	lake	sal	0.05	734	DR-3	1939-01-07	DR-3	-47.15	
6	734	pb	temp	-21.50	734	DR-3	1939-01-07	DR-3	-47.15	
7	735	pb	rad	7.22	735	DR-3	1930-01-12	DR-3	-47.15	
8	735	NaN	sal	0.06	735	DR-3	1930-01-12	DR-3	-47.15	
9	735	NaN	temp	-26.00	735	DR-3	1930-01-12	DR-3	-47.15	
10	751	pb	rad	4.35	751	DR-3	1930-02-26	DR-3	-47.15	
11	751	pb	temp	-18.50	751	DR-3	1930-02-26	DR-3	-47.15	
12	751	lake	sal	0.10	751	DR-3	1930-02-26	DR-3	-47.15	
13	752	lake	rad	2.19	752	DR-3	NaN	DR-3	-47.15	
14	752	lake	sal	0.09	752	DR-3	NaN	DR-3	-47.15	
15	752	lake	temp	-16.00	752	DR-3	NaN	DR-3	-47.15	
16	752	roe	sal	41.60	752	DR-3	NaN	DR-3	-47.15	
17	837	lake	rad	1.46	837	MSK-4	1932-01-14	MSK-4	-48.87	
18	837	lake	sal	0.21	837	MSK-4	1932-01-14	MSK-4	-48.87	
19	837	roe	sal	22.50	837	MSK-4	1932-01-14	MSK-4	-48.87	
20	844	roe	rad	11.25	844	DR-1	1932-03-22	DR-1	-49.85	

	long
0	-128.57
1	-128.57
2	-128.57
3	-128.57
4	-126.72
5	-126.72
6	-126.72
7	-126.72
8	-126.72
9	-126.72
10	-126.72

```

11 -126.72
12 -126.72
13 -126.72
14 -126.72
15 -126.72
16 -126.72
17 -123.40
18 -123.40
19 -123.40
20 -128.57

```

Lastly, we add the person information to this dataset.

```
[106]: merged = s2v2loc_merge.merge(person, how = 'left', left_on = 'person', right_on=
      ↪='ident')
      print(merged.head())
```

	taken	person	quant	reading	ident_x	site	dated	name	lat	\
0	619	dye	rad	9.82	619	DR-1	1927-02-08	DR-1	-49.85	
1	619	dye	sal	0.13	619	DR-1	1927-02-08	DR-1	-49.85	
2	622	dye	rad	7.80	622	DR-1	1927-02-10	DR-1	-49.85	
3	622	dye	sal	0.09	622	DR-1	1927-02-10	DR-1	-49.85	
4	734	pb	rad	8.41	734	DR-3	1939-01-07	DR-3	-47.15	

	long	ident_y	personal	family
0	-128.57	dye	William	Dye
1	-128.57	dye	William	Dye
2	-128.57	dye	William	Dye
3	-128.57	dye	William	Dye
4	-126.72	pb	Frank	Pabodie

You can merge based on multiple columns as long as they match up.

```
[107]: ps = person.merge(survey, left_on = 'ident', right_on = 'person')
      vs = visited.merge(survey, left_on = 'ident', right_on = 'taken')
      print(ps)
```

	ident	personal	family	taken	person	quant	reading
0	dye	William	Dye	619	dye	rad	9.82
1	dye	William	Dye	619	dye	sal	0.13
2	dye	William	Dye	622	dye	rad	7.80
3	dye	William	Dye	622	dye	sal	0.09
4	pb	Frank	Pabodie	734	pb	rad	8.41
5	pb	Frank	Pabodie	734	pb	temp	-21.50
6	pb	Frank	Pabodie	735	pb	rad	7.22
7	pb	Frank	Pabodie	751	pb	rad	4.35
8	pb	Frank	Pabodie	751	pb	temp	-18.50
9	lake	Anderson	Lake	734	lake	sal	0.05
10	lake	Anderson	Lake	751	lake	sal	0.10

11	lake	Anderson	Lake	752	lake	rad	2.19
12	lake	Anderson	Lake	752	lake	sal	0.09
13	lake	Anderson	Lake	752	lake	temp	-16.00
14	lake	Anderson	Lake	837	lake	rad	1.46
15	lake	Anderson	Lake	837	lake	sal	0.21
16	roe	Valentina	Roerich	752	roe	sal	41.60
17	roe	Valentina	Roerich	837	roe	sal	22.50
18	roe	Valentina	Roerich	844	roe	rad	11.25

```
[108]: print(vs)
```

	ident	site	dated	taken	person	quant	reading
0	619	DR-1	1927-02-08	619	dye	rad	9.82
1	619	DR-1	1927-02-08	619	dye	sal	0.13
2	622	DR-1	1927-02-10	622	dye	rad	7.80
3	622	DR-1	1927-02-10	622	dye	sal	0.09
4	734	DR-3	1939-01-07	734	pb	rad	8.41
5	734	DR-3	1939-01-07	734	lake	sal	0.05
6	734	DR-3	1939-01-07	734	pb	temp	-21.50
7	735	DR-3	1930-01-12	735	pb	rad	7.22
8	735	DR-3	1930-01-12	735	NaN	sal	0.06
9	735	DR-3	1930-01-12	735	NaN	temp	-26.00
10	751	DR-3	1930-02-26	751	pb	rad	4.35
11	751	DR-3	1930-02-26	751	pb	temp	-18.50
12	751	DR-3	1930-02-26	751	lake	sal	0.10
13	752	DR-3	NaN	752	lake	rad	2.19
14	752	DR-3	NaN	752	lake	sal	0.09
15	752	DR-3	NaN	752	lake	temp	-16.00
16	752	DR-3	NaN	752	roe	sal	41.60
17	837	MSK-4	1932-01-14	837	lake	rad	1.46
18	837	MSK-4	1932-01-14	837	lake	sal	0.21
19	837	MSK-4	1932-01-14	837	roe	sal	22.50
20	844	DR-1	1932-03-22	844	roe	rad	11.25

```
[109]: ps_vs = ps.merge(vs,
                        left_on = ['ident', 'taken', 'quant', 'reading'],
                        right_on = ['person', 'ident', 'quant', 'reading']) # The keys
                        ↪ need to correspond
ps_vs.head()
```

```
[109]:  ident_x personal  family  taken_x person_x quant  reading  ident_y  site \
0      dye William    Dyer    619      dye    rad    9.82    619 DR-1
1      dye William    Dyer    619      dye    sal    0.13    619 DR-1
2      dye William    Dyer    622      dye    rad    7.80    622 DR-1
3      dye William    Dyer    622      dye    sal    0.09    622 DR-1
4      pb  Frank  Pabodie    734      pb    rad    8.41    734 DR-3
```


	dated	taken_y	person_y
0	1927-02-08	619	dyer
1	1927-02-08	619	dyer
2	1927-02-10	622	dyer
3	1927-02-10	622	dyer
4	1939-01-07	734	pb

Note that since there are common column names, the merge appends `_x` and `_y` to denote which column came from the left and right, respectively.

1.6.4 Tidy data principles and reshaping datasets

The tidy data principle is a principle espoused by Dr. Hadley Wickham, one of the foremost R developers. [Tidy data](#) is a structure for datasets to make them more easily analyzed on computers. The basic principles are

- Each row is an observation
- Each column is a variable
- Each type of observational unit forms a table

Tidy data is tidy in one way. Untidy data can be untidy in many ways

Let's look at some examples.

```
[110]: from glob import glob
        filenames = sorted(glob('data/table*.csv')) # find files matching pattern. I
        ↪ know there are 6 of them
        table1, table2, table3, table4a, table4b, table5 = [pd.read_csv(f) for f in
        ↪ filenames] # Use a list comprehension
```

This code imports data from 6 files matching a pattern. Python allows multiple assignments on the left of the `=`, and as each dataset is imported, it gets assigned in order to the variables on the left. In the second line I sort the file names so that they match the order in which I'm storing them in the 3rd line. The function `glob` does pattern-matching of file names.

The following tables refer to the number of TB cases and population in Afghanistan, Brazil and China in 1999 and 2000

```
[111]: print(table1)
```

	country	year	cases	population
0	Afghanistan	1999	745	19987071
1	Afghanistan	2000	2666	20595360
2	Brazil	1999	37737	172006362
3	Brazil	2000	80488	174504898
4	China	1999	212258	1272915272
5	China	2000	213766	1280428583

```
[112]: print(table2)
```

	country	year	type	count
0	Afghanistan	1999	cases	745
1	Afghanistan	1999	population	19987071
2	Afghanistan	2000	cases	2666
3	Afghanistan	2000	population	20595360
4	Brazil	1999	cases	37737
5	Brazil	1999	population	172006362
6	Brazil	2000	cases	80488
7	Brazil	2000	population	174504898
8	China	1999	cases	212258
9	China	1999	population	1272915272
10	China	2000	cases	213766
11	China	2000	population	1280428583

```
[113]: print(table3)
```

	country	year	rate
0	Afghanistan	1999	745/19987071
1	Afghanistan	2000	2666/20595360
2	Brazil	1999	37737/172006362
3	Brazil	2000	80488/174504898
4	China	1999	212258/1272915272
5	China	2000	213766/1280428583

```
[114]: print(table4a) # cases
```

	country	1999	2000
0	Afghanistan	745	2666
1	Brazil	37737	80488
2	China	212258	213766

```
[115]: print(table4b) # population
```

	country	1999	2000
0	Afghanistan	19987071	20595360
1	Brazil	172006362	174504898
2	China	1272915272	1280428583

```
[116]: print(table5)
```

	country	century	year	rate
0	Afghanistan	19	99	745/19987071
1	Afghanistan	20	0	2666/20595360
2	Brazil	19	99	37737/172006362
3	Brazil	20	0	80488/174504898
4	China	19	99	212258/1272915272
5	China	20	0	213766/1280428583

Exercise: Describe why and why not each of these datasets are tidy.

1.6.5 Melting (unpivoting) data

Melting is the operation of collapsing multiple columns into 2 columns, where one column is formed by the old column names, and the other by the corresponding values. Some columns may be kept fixed and their data are repeated to maintain the interrelationships between the variables.

We'll start with loading some data on income and religion in the US from the Pew Research Center.

```
[117]: pew = pd.read_csv('data/pew.csv')
print(pew.head())
```

	religion	<\$10k	\$10-20k	\$20-30k	\$30-40k	\$40-50k	\$50-75k	\
0	Agnostic	27	34	60	81	76	137	
1	Atheist	12	27	37	52	35	70	
2	Buddhist	27	21	30	34	33	58	
3	Catholic	418	617	732	670	638	1116	
4	Don't know/refused	15	14	15	11	10	35	

	\$75-100k	\$100-150k	>150k	Don't know/refused
0	122	109	84	96
1	73	59	74	76
2	62	39	53	54
3	949	792	633	1489
4	21	17	18	116

This dataset is considered in “wide” format. There are several issues with it, including the fact that column headers have data. Those column headers are income groups, that should be a column by tidy principles. Our job is to turn this dataset into “long” format with a column for income group.

We will use the function `melt` to achieve this. This takes a few parameters:

- **id_vars** is a list of variables that will remain as is
- **value_vars** is a list of column names that we will melt (or unpivot). By default, it will melt all columns not mentioned in `id_vars`
- **var_name** is a string giving the name of the new column created by the headers (default: `variable`)
- **value_name** is a string giving the name of the new column created by the values (default: `value`)

```
[118]: pew_long = pew.melt(id_vars = ['religion'], var_name = 'income_group',
    ↪ value_name = 'count')
print(pew_long.head())
```

	religion	income_group	count
0	Agnostic	<\$10k	27
1	Atheist	<\$10k	12
2	Buddhist	<\$10k	27
3	Catholic	<\$10k	418
4	Don't know/refused	<\$10k	15

1.6.6 Separating columns containing multiple variables

We will use an Ebola dataset to illustrate this principle

```
[119]: ebola = pd.read_csv('data/country_timeseries.csv')
print(ebola.head())
```

	Date	Day	Cases_Guinea	Cases_Liberia	Cases_SierraLeone	\
0	1/5/2015	289	2776.0	NaN	10030.0	
1	1/4/2015	288	2775.0	NaN	9780.0	
2	1/3/2015	287	2769.0	8166.0	9722.0	
3	1/2/2015	286	NaN	8157.0	NaN	
4	12/31/2014	284	2730.0	8115.0	9633.0	

	Cases_Nigeria	Cases_Senegal	Cases_UnitedStates	Cases_Spain	Cases_Mali	\
0	NaN	NaN	NaN	NaN	NaN	
1	NaN	NaN	NaN	NaN	NaN	
2	NaN	NaN	NaN	NaN	NaN	
3	NaN	NaN	NaN	NaN	NaN	
4	NaN	NaN	NaN	NaN	NaN	

	Deaths_Guinea	Deaths_Liberia	Deaths_SierraLeone	Deaths_Nigeria	\
0	1786.0	NaN	2977.0	NaN	
1	1781.0	NaN	2943.0	NaN	
2	1767.0	3496.0	2915.0	NaN	
3	NaN	3496.0	NaN	NaN	
4	1739.0	3471.0	2827.0	NaN	

	Deaths_Senegal	Deaths_UnitedStates	Deaths_Spain	Deaths_Mali
0	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN

Note that for each country we have two columns – one for cases (number infected) and one for deaths. Ideally we want one column for country, one for cases and one for deaths.

The first step will be to melt this data sets so that the column headers in question from a column and the corresponding data forms a second column.

```
[120]: ebola_long = ebola.melt(id_vars = ['Date', 'Day'])
print(ebola_long.head())
```

	Date	Day	variable	value
0	1/5/2015	289	Cases_Guinea	2776.0
1	1/4/2015	288	Cases_Guinea	2775.0
2	1/3/2015	287	Cases_Guinea	2769.0
3	1/2/2015	286	Cases_Guinea	NaN
4	12/31/2014	284	Cases_Guinea	2730.0

We now need to split the data in the `variable` column to make two columns. One will contain the country name and the other either Cases or Deaths. We will use some string manipulation functions that we will see later to achieve this.

```
[121]: variable_split = ebola_long['variable'].str.split('_', expand=True) # split on
      ↳ the `_` character
      print(variable_split[:5])
```

```

      0      1
0  Cases  Guinea
1  Cases  Guinea
2  Cases  Guinea
3  Cases  Guinea
4  Cases  Guinea
```

The `expand=True` option forces the creation of an `DataFrame` rather than a list

```
[122]: type(variable_split)
```

```
[122]: pandas.core.frame.DataFrame
```

We can now concatenate this to the original data

```
[123]: variable_split.columns = ['status', 'country']

      ebola_parsed = pd.concat([ebola_long, variable_split], axis = 1)

      ebola_parsed.drop('variable', axis = 1, inplace=True) # Remove the column named
      ↳ "variable" and replace the old data with the new one in the same location

      print(ebola_parsed.head())
```

```

      Date  Day  value status country
0  1/5/2015  289  2776.0  Cases  Guinea
1  1/4/2015  288  2775.0  Cases  Guinea
2  1/3/2015  287  2769.0  Cases  Guinea
3  1/2/2015  286     NaN  Cases  Guinea
4  12/31/2014  284  2730.0  Cases  Guinea
```

1.6.7 Pivot/spread datasets

If we wanted to, we could also make two columns based on cases and deaths, so for each country and date you could easily read off the cases and deaths. This is achieved using the `pivot_table` function.

In the `pivot_table` syntax, `index` refers to the columns we don't want to change, `columns` refers to the column whose values will form the column names of the new columns, and `values` is the name of the column that will form the values in the pivoted dataset.

```
[124]: ebola_parsed.pivot_table(index = ['Date', 'Day', 'country'], columns = 'status',
    ↪values = 'value')
```

```
[124]: status      Cases  Deaths
Date    Day country
1/2/2015 286 Liberia    8157.0  3496.0
1/3/2015 287 Guinea    2769.0  1767.0
          Liberia    8166.0  3496.0
          SierraLeone 9722.0  2915.0
1/4/2015 288 Guinea    2775.0  1781.0
...
9/7/2014 169 Liberia    2081.0  1137.0
          Nigeria      21.0     8.0
          Senegal       3.0     0.0
          SierraLeone 1424.0  524.0
9/9/2014 171 Liberia    2407.0   NaN
```

[375 rows x 2 columns]

This creates something called MultiIndex in the pandas DataFrame. This is useful in some advanced cases, but here, we just want a normal DataFrame back. We can achieve that by using the `reset_index` function.

```
[125]: ebola_parsed.pivot_table(index = ['Date', 'Day', 'country'], columns = 'status',
    ↪values = 'value').reset_index()
```

```
[125]: status      Date  Day    country  Cases  Deaths
0      1/2/2015  286    Liberia    8157.0  3496.0
1      1/3/2015  287     Guinea    2769.0  1767.0
2      1/3/2015  287    Liberia    8166.0  3496.0
3      1/3/2015  287  SierraLeone  9722.0  2915.0
4      1/4/2015  288     Guinea    2775.0  1781.0
..      ...    ...
370    9/7/2014  169    Liberia    2081.0  1137.0
371    9/7/2014  169    Nigeria      21.0     8.0
372    9/7/2014  169    Senegal       3.0     0.0
373    9/7/2014  169  SierraLeone  1424.0  524.0
374    9/9/2014  171    Liberia    2407.0   NaN
```

[375 rows x 5 columns]

Pivoting is a 2-column to many-column operation, with the number of columns formed depending on the number of unique values present in the column of the original data that is entered into the `columns` argument of `pivot_table`

Exercise: Load the file `weather.csv` into Python and work on making it a tidy dataset. It requires melting and pivoting. The dataset comprises of the maximum and minimum temperatures recorded each day in 2010. There are lots of missing value. Ultimately we want columns for days of the

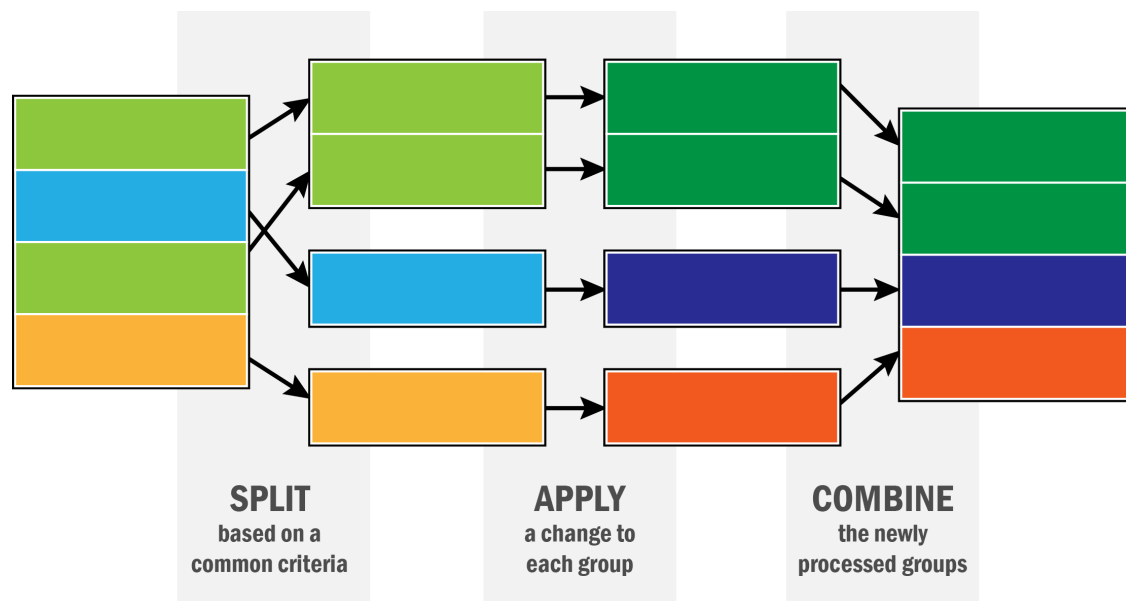
month, maximum temperature and minimum temperature along with the location ID, the year and the month.

1.7 Data aggregation and split-apply-combine

We'll use the Gapminder dataset for this section

```
[126]: df = pd.read_csv('data/gapminder.tsv', sep = '\t') # data is tab-separated, so
        ↪ we use '\t' to specify that
```

The paradigm we will be exploring is often called *split-apply-combine* or MapReduce or grouped aggregation. The basic idea is that you split a data set up by some feature, apply a recipe to each piece, compute the result, and then put the results back together into a dataset. This can be described in the following schematic.



pandas is set up for this. It features the `groupby` function that allows the “split” part of the operation. We can then apply a function to each part and put it back together. Let's see how.

```
[127]: df.head()
```

```
[127]:
```

	country	continent	year	lifeExp	pop	gdpPercap
0	Afghanistan	Asia	1952	28.801	8425333	779.445314
1	Afghanistan	Asia	1957	30.332	9240934	820.853030
2	Afghanistan	Asia	1962	31.997	10267083	853.100710
3	Afghanistan	Asia	1967	34.020	11537966	836.197138
4	Afghanistan	Asia	1972	36.088	13079460	739.981106

```
[128]: f"This dataset has {len(df['country'].unique())} countries in it"
```

```
[128]: 'This dataset has 142 countries in it'
```

One of the variables in this dataset is life expectancy at birth, `lifeExp`. Suppose we want to find the average life expectancy of each country over the period of study.

```
[129]: df.groupby('country')['lifeExp'].mean()
```

```
[129]: country
Afghanistan      37.478833
Albania          68.432917
Algeria          59.030167
Angola           37.883500
Argentina        69.060417
...
Vietnam          57.479500
West Bank and Gaza 60.328667
Yemen, Rep.      46.780417
Zambia           45.996333
Zimbabwe         52.663167
Name: lifeExp, Length: 142, dtype: float64
```

So what's going on here? First, we use the `groupby` function, telling pandas to split the dataset up by values of the column `country`.

```
[130]: df.groupby('country')
```

```
[130]: <pandas.core.groupby.generic.DataFrameGroupBy object at 0x7feba93dc590>
```

pandas won't show you the actual data, but will tell you that it is a grouped dataframe object. This means that each element of this object is a `DataFrame` with data from one country.

```
[131]: df.groupby('country').ngroups
```

```
[131]: 142
```

```
[132]: df.groupby('country').get_group('United Kingdom')
```

```
[132]:
```

	country	continent	year	lifeExp	pop	gdpPercap
1596	United Kingdom	Europe	1952	69.180	50430000	9979.508487
1597	United Kingdom	Europe	1957	70.420	51430000	11283.177950
1598	United Kingdom	Europe	1962	70.760	53292000	12477.177070
1599	United Kingdom	Europe	1967	71.360	54959000	14142.850890
1600	United Kingdom	Europe	1972	72.010	56079000	15895.116410
1601	United Kingdom	Europe	1977	72.760	56179000	17428.748460
1602	United Kingdom	Europe	1982	74.040	56339704	18232.424520
1603	United Kingdom	Europe	1987	75.007	56981620	21664.787670
1604	United Kingdom	Europe	1992	76.420	57866349	22705.092540
1605	United Kingdom	Europe	1997	77.218	58808266	26074.531360
1606	United Kingdom	Europe	2002	78.471	59912431	29478.999190
1607	United Kingdom	Europe	2007	79.425	60776238	33203.261280


```
[133]: type(df.groupby('country').get_group('United Kingdom'))
```

```
[133]: pandas.core.frame.DataFrame
```

```
[134]: avg_lifeexp_country = df.groupby('country').lifeExp.mean()  
avg_lifeexp_country['United Kingdom']
```

```
[134]: 73.92258333333332
```

```
[135]: df.groupby('country').get_group('United Kingdom').lifeExp.mean()
```

```
[135]: 73.92258333333332
```

Let's look at if life expectancy has gone up over time, by continent

```
[136]: df.groupby(['continent', 'year']).lifeExp.mean()
```

```
[136]: continent  year  
Africa      1952    39.135500  
            1957    41.266346  
            1962    43.319442  
            1967    45.334538  
            1972    47.450942  
            1977    49.580423  
            1982    51.592865  
            1987    53.344788  
            1992    53.629577  
            1997    53.598269  
            2002    53.325231  
            2007    54.806038  
Americas    1952    53.279840  
            1957    55.960280  
            1962    58.398760  
            1967    60.410920  
            1972    62.394920  
            1977    64.391560  
            1982    66.228840  
            1987    68.090720  
            1992    69.568360  
            1997    71.150480  
            2002    72.422040  
            2007    73.608120  
Asia        1952    46.314394  
            1957    49.318544  
            1962    51.563223  
            1967    54.663640  
            1972    57.319269
```

	1977	59.610556
	1982	62.617939
	1987	64.851182
	1992	66.537212
	1997	68.020515
	2002	69.233879
	2007	70.728485
Europe	1952	64.408500
	1957	66.703067
	1962	68.539233
	1967	69.737600
	1972	70.775033
	1977	71.937767
	1982	72.806400
	1987	73.642167
	1992	74.440100
	1997	75.505167
	2002	76.700600
	2007	77.648600
Oceania	1952	69.255000
	1957	70.295000
	1962	71.085000
	1967	71.310000
	1972	71.910000
	1977	72.855000
	1982	74.290000
	1987	75.320000
	1992	76.945000
	1997	78.190000
	2002	79.740000
	2007	80.719500

Name: lifeExp, dtype: float64

```
[137]: avg_lifeexp_continent_yr = df.groupby(['continent','year']).lifeExp.mean().
        ↪reset_index()
        avg_lifeexp_continent_yr
```

```
[137]:   continent  year  lifeExp
0     Africa  1952  39.135500
1     Africa  1957  41.266346
2     Africa  1962  43.319442
3     Africa  1967  45.334538
4     Africa  1972  47.450942
5     Africa  1977  49.580423
6     Africa  1982  51.592865
7     Africa  1987  53.344788
8     Africa  1992  53.629577
```

9	Africa	1997	53.598269
10	Africa	2002	53.325231
11	Africa	2007	54.806038
12	Americas	1952	53.279840
13	Americas	1957	55.960280
14	Americas	1962	58.398760
15	Americas	1967	60.410920
16	Americas	1972	62.394920
17	Americas	1977	64.391560
18	Americas	1982	66.228840
19	Americas	1987	68.090720
20	Americas	1992	69.568360
21	Americas	1997	71.150480
22	Americas	2002	72.422040
23	Americas	2007	73.608120
24	Asia	1952	46.314394
25	Asia	1957	49.318544
26	Asia	1962	51.563223
27	Asia	1967	54.663640
28	Asia	1972	57.319269
29	Asia	1977	59.610556
30	Asia	1982	62.617939
31	Asia	1987	64.851182
32	Asia	1992	66.537212
33	Asia	1997	68.020515
34	Asia	2002	69.233879
35	Asia	2007	70.728485
36	Europe	1952	64.408500
37	Europe	1957	66.703067
38	Europe	1962	68.539233
39	Europe	1967	69.737600
40	Europe	1972	70.775033
41	Europe	1977	71.937767
42	Europe	1982	72.806400
43	Europe	1987	73.642167
44	Europe	1992	74.440100
45	Europe	1997	75.505167
46	Europe	2002	76.700600
47	Europe	2007	77.648600
48	Oceania	1952	69.255000
49	Oceania	1957	70.295000
50	Oceania	1962	71.085000
51	Oceania	1967	71.310000
52	Oceania	1972	71.910000
53	Oceania	1977	72.855000
54	Oceania	1982	74.290000
55	Oceania	1987	75.320000

```

56 Oceania 1992 76.945000
57 Oceania 1997 78.190000
58 Oceania 2002 79.740000
59 Oceania 2007 80.719500

```

```
[138]: type(avg_lifeexp_continent_yr)
```

```
[138]: pandas.core.frame.DataFrame
```

The aggregation function, in this case `mean`, does both the “apply” and “combine” parts of the process.

We can do quick aggregations with `pandas`

```
[139]: df.groupby('continent').lifeExp.describe()
```

```
[139]:
```

	count	mean	std	min	25%	50%	75%	\
continent								
Africa	624.0	48.865330	9.150210	23.599	42.37250	47.7920	54.41150	
Americas	300.0	64.658737	9.345088	37.579	58.41000	67.0480	71.69950	
Asia	396.0	60.064903	11.864532	28.801	51.42625	61.7915	69.50525	
Europe	360.0	71.903686	5.433178	43.585	69.57000	72.2410	75.45050	
Oceania	24.0	74.326208	3.795611	69.120	71.20500	73.6650	77.55250	


```

max
continent
Africa      76.442
Americas    80.653
Asia        82.603
Europe      81.757
Oceania     81.235

```

```
[140]: df.groupby('continent').nth(10) # Tenth observation in each group
```

```
[140]:
```

	country	year	lifeExp	pop	gdpPercap
continent					
Africa	Algeria	2002	70.994	31287142	5288.040382
Americas	Argentina	2002	74.340	38331121	8797.640716
Asia	Afghanistan	2002	42.129	25268405	726.734055
Europe	Albania	2002	75.651	3508512	4604.211737
Oceania	Australia	2002	80.370	19546792	30687.754730

You can also use functions from other modules, or your own functions in this aggregation work.

```
[141]: df.groupby('continent').lifeExp.agg(np.mean)
```

```
[141]: continent
Africa      48.865330
```

```
Americas    64.658737
Asia        60.064903
Europe      71.903686
Oceania     74.326208
Name: lifeExp, dtype: float64
```

```
[142]: def my_mean(values):
        n = len(values)
        sum = 0
        for value in values:
            sum += value
        return(sum/n)

df.groupby('continent').lifeExp.agg(my_mean)
```

```
[142]: continent
Africa    48.865330
Americas  64.658737
Asia      60.064903
Europe    71.903686
Oceania   74.326208
Name: lifeExp, dtype: float64
```

You can do many functions at once

```
[143]: df.groupby('year').lifeExp.agg([np.count_nonzero, np.mean, np.std])
```

```
[143]:      count_nonzero      mean      std
year
1952           142.0  49.057620  12.225956
1957           142.0  51.507401  12.231286
1962           142.0  53.609249  12.097245
1967           142.0  55.678290  11.718858
1972           142.0  57.647386  11.381953
1977           142.0  59.570157  11.227229
1982           142.0  61.533197  10.770618
1987           142.0  63.212613  10.556285
1992           142.0  64.160338  11.227380
1997           142.0  65.014676  11.559439
2002           142.0  65.694923  12.279823
2007           142.0  67.007423  12.073021
```

You can also aggregate on different columns at the same time by passing a dict to the agg function

```
[144]: df.groupby('year').agg({'lifeExp': np.mean, 'pop': np.median, 'gdpPercap': np.
    ↳median}).reset_index()
```

```
[144]:   year    lifeExp      pop  gdpPercap
      0  1952  49.057620  3943953.0  1968.528344
      1  1957  51.507401  4282942.0  2173.220291
      2  1962  53.609249  4686039.5  2335.439533
      3  1967  55.678290  5170175.5  2678.334741
      4  1972  57.647386  5877996.5  3339.129407
      5  1977  59.570157  6404036.5  3798.609244
      6  1982  61.533197  7007320.0  4216.228428
      7  1987  63.212613  7774861.5  4280.300366
      8  1992  64.160338  8688686.5  4386.085502
      9  1997  65.014676  9735063.5  4781.825478
     10  2002  65.694923 10372918.5  5319.804524
     11  2007  67.007423 10517531.0  6124.371109
```

Transformation You can do grouped transformations using this same method. We will compute the z-score for each year, i.e. we will subtract the average life expectancy and divide by the standard deviation

```
[145]: def my_zscore(values):
      m = np.mean(values)
      s = np.std(values)
      return((values - m)/s)
```

```
[146]: df.groupby('year').lifeExp.transform(my_zscore)
```

```
[146]: 0      -1.662719
      1      -1.737377
      2      -1.792867
      3      -1.854699
      4      -1.900878
      ...
     1699    -0.081910
     1700    -0.338167
     1701    -1.580537
     1702    -2.100756
     1703    -1.955077
      Name: lifeExp, Length: 1704, dtype: float64
```

```
[147]: df['lifeExp_z'] = df.groupby('year').lifeExp.transform(my_zscore)
```

```
[148]: df.groupby('year').lifeExp_z.mean()
```

```
[148]: year
      1952   -1.103089e-15
      1957    1.802842e-15
      1962    1.464400e-15
      1967   -1.935072e-17
```

```

1972    -1.057448e-15
1977     2.447182e-16
1982     1.122928e-15
1987    -2.045899e-15
1992     5.773942e-16
1997    -1.590277e-15
2002     5.254013e-16
2007     5.035096e-16
Name: lifeExp_z, dtype: float64

```

Filter We can split the dataset by values of one variable, and filter out those splits that fail some criterion. The following code only keeps countries with a population of at least 10 million at some point during the study period

```
[149]: df.groupby('country').filter(lambda d: d['pop'].max() > 10000000)
```

```
[149]:
```

	country	continent	year	lifeExp	pop	gdpPercap	lifeExp_z
0	Afghanistan	Asia	1952	28.801	8425333	779.445314	-1.662719
1	Afghanistan	Asia	1957	30.332	9240934	820.853030	-1.737377
2	Afghanistan	Asia	1962	31.997	10267083	853.100710	-1.792867
3	Afghanistan	Asia	1967	34.020	11537966	836.197138	-1.854699
4	Afghanistan	Asia	1972	36.088	13079460	739.981106	-1.900878
...
1699	Zimbabwe	Africa	1987	62.351	9216418	706.157306	-0.081910
1700	Zimbabwe	Africa	1992	60.377	10704340	693.420786	-0.338167
1701	Zimbabwe	Africa	1997	46.809	11404948	792.449960	-1.580537
1702	Zimbabwe	Africa	2002	39.989	11926563	672.038623	-2.100756
1703	Zimbabwe	Africa	2007	43.487	12311143	469.709298	-1.955077

```
[924 rows x 7 columns]
```